

CS 143 Lab 5 Program Specifications

Checkpoint Due 11:59PM Friday Aug 10

Lab Due 11:59PM Monday Aug 13

Problem Overview: This assignment will give you more practice with recursion, and exposure to the important concepts of *regular expressions* and *grammars*. You will also get the opportunity to work with the `JCF Map` and `Set` classes. These outcomes will be achieved when you use a client program that reads an input file with a grammar in *Backus-Naur Form* (BNF) and allows the user to randomly generate elements of the grammar using the `GrammarProcessor` class that you will implement.

The `Lab6Files` folder posted with this assignment contains two client programs

`GrammarTest.java` and `GrammarProcessorClient.java`.

- These client files perform the file processing and user interaction for your program.
- `GrammarTest` is a simple tester for the `GrammarProcessor` class that you will write. It may be easier for you to use and modify as needed in the early stages of implementing your `GrammarProcessor` class.
- `GrammarProcessorClient` is more versatile and user-friendly for testing your `GrammarProcessor` class thoroughly in the final stages of development.

BNF Grammars: A grammar will be specified as a sequence of Strings, each of which represents the rules for a nonterminal symbol. Each String will be of the form:

```
<nonterminal symbol> ::= <rule> | <rule> | <rule> | ... | <rule>
```

In the online *Introduction to Programming in Java 7/e* textbook by David J. Eck, Section 9.5.1 presented the following such strings:

```
<sentence> ::= <noun-phrase> | <verb-phrase>
<verb-phrase> ::= <intransitive-verb> | <transitive-verb> <noun-phrase>
```

In practice, the grammar may or may not use `<>` characters, and for our program, since it will be asking users to type nonterminal symbols from a BNF grammar, the symbols will be very short as in BNF grammar file `sentence1.txt` shown here:

```
<s> ::= <np> <vp>
<np> ::= <dp> <adjp> <n> | <pn>
<pn> ::= John | Jane | Sally | Spot | Fred | Elmo
<adjp> ::= <adj> | <adj> <adjp>
<adj> ::= big | fat | green | wonderful | faulty | subliminal | pretentious
<dp> ::= the | a
<n> ::= dog | cat | man | university | father | mother | child | television
<vp> ::= <tv> <np> | <iv>
<tv> ::= hit | honored | kissed | helped
<iv> ::= died | collapsed | laughed | wept
```

Here is the BNF grammar contained in `sentence2.txt` (notice it doesn't use `<>` around its symbols):

```
E ::= T | E OP T
T ::= x | y | 42 | 0 | 1 | 92 | ( E ) | F1 ( E ) | - T | F2 ( E , E )
OP ::= + | - | * | % | /
F1 ::= sin | cos | tan | sqrt | abs
F2 ::= max | min | pow
```

Sample I/O sessions shown below are based on a working `GrammarProcessor` class and the `GrammarProcessorClient`. User input is shown underlined only for easy identification. Since the language components are randomly generated, you should not expect your program to produce the exact same I/O sessions if the user types the exact same input values.

Sample I/O Session for Sentence 1 Grammar:

What is the name of the grammar file? sentencel.txt

[<adj>, <adjp>, <dp>, <iv>, <n>, <np>, <pn>, <s>, <tv>, <vp>]

Which symbol do you want? (press <enter> to quit) <s>

How many expressions? 5

the subliminal cat kissed John

Elmo honored Sally

a fat subliminal man hit Jane

the big subliminal fat cat collapsed

Jane collapsed

[<adj>, <adjp>, <dp>, <iv>, <n>, <np>, <pn>, <s>, <tv>, <vp>]

Which symbol do you want? (press <enter> to quit)

Sample I/O Session for Sentence 2 Grammar:

[E, F1, F2, OP, T]

Which symbol do you want? (press <enter> to quit) E

How many expressions? 5

x

92

min (x * 0 , - 1)

1 - x * 0

- - 1 + (y % x - 42)

[E, F1, F2, OP, T]

Which symbol do you want? (press <enter> to quit)

Regular Expressions and Useful String Methods: Be sure to study the `ReadMary.java` program posted with this assignment. It shows how *regular expressions* can be used to set the delimiters on a `Scanner` object to allow characters other than just whitespace to be ignored while reading tokens from the input stream connected to the `Scanner`.

Notice that each line in the BNF grammars shown above contains exactly one occurrence of "`: :=`" to separate the nonterminal symbol on the left-hand side from the rules on the right-hand side. Any token that appears to the left of "`: :=`" in the grammar is considered a nonterminal. All other tokens are considered terminals. The right-hand side will contain the "or" symbol ("`|`") to separate one rule from another, unless there is only one rule for that nonterminal symbol. Each of the rules will have a series of tokens (always at least one) separated and potentially surrounded by whitespace. There could be any amount of whitespace surrounding tokens.

The `String` class has the following method: `String[] split(String regularExpression)`
It can be used to split a string into substrings delimited by the regular expression. Suppose that `line` is a string with the following value: "`<ex>: :=yours|mine|ours`"

Then this statement

```
String[] parts = line.split(": :=");
```

has the effect of creating `parts[0]` containing "<ex>" and `parts[1]` containing "yours|mine|ours".

What if `line` has the value "yours|mine|ours"? In order to split the line into substrings delimited by "|", you must use the following syntax:

```
String[] parts = line.split("[|]");
```

That is, you must enclose the "|" symbol inside [] because the "|" symbol has special meaning in regular expressions. The preceding statement results in `parts[0]` holding "yours", `parts[1]` holding "mine", and `parts[2]` holding "ours".

Lastly, suppose that `line` has the following value: " to be or not to be"

There are whitespace characters separating the symbols in each rule from each other. To split the string into substrings delimited by whitespace, we would use the following statement:

```
String[] parts = line.split("[ \\t]+");
```

This says there are one or more spaces or tabs to ignore between splits. The only issue with the statement above is that if `line` begins with a whitespace character, `parts[0]` will contain the empty string "".

So another `String` class method that may be useful for this assignment is the following:

`String trim()` can be used to create a string that has leading and trailing spaces removed.

```
String st = s.trim(); //st holds trimmed version of s; s is unchanged.
```

Java Collections Framework: You will be using `Set`, `Map` and possibly `Iterator` objects in this assignment. Those objects will handle all your data storage needs; you just need to know how to create and access them. Specifically, you will be storing grammar rules in the following type of object:

`TreeMap<String, String[]>` object.

Think about it... Draw a picture of sample grammar rules stored in a map so that you understand why this choice of data structure makes sense for this lab. Make sure that you understand what the `get()`, `put()`, `containsKey()`, `keyset()` map methods listed in Table 11.5 on page 726 (704 in 2/e) do. Since the `keyset()` method returns a `Set` object, be sure that you understand how to traverse a `Set` using either an `Iterator` object, or a `for-each` loop.

GrammarProcessor Specifications: `GrammarTest` and `GrammarProcessorClient` will read a BNF grammar from a text file and store each line from the file into a `List<String>` that is passed to your `GrammarProcessor` constructor. Your constructor will store the grammar in a manner that makes it convenient to generate random elements of the grammar, as shown in the specification table below.

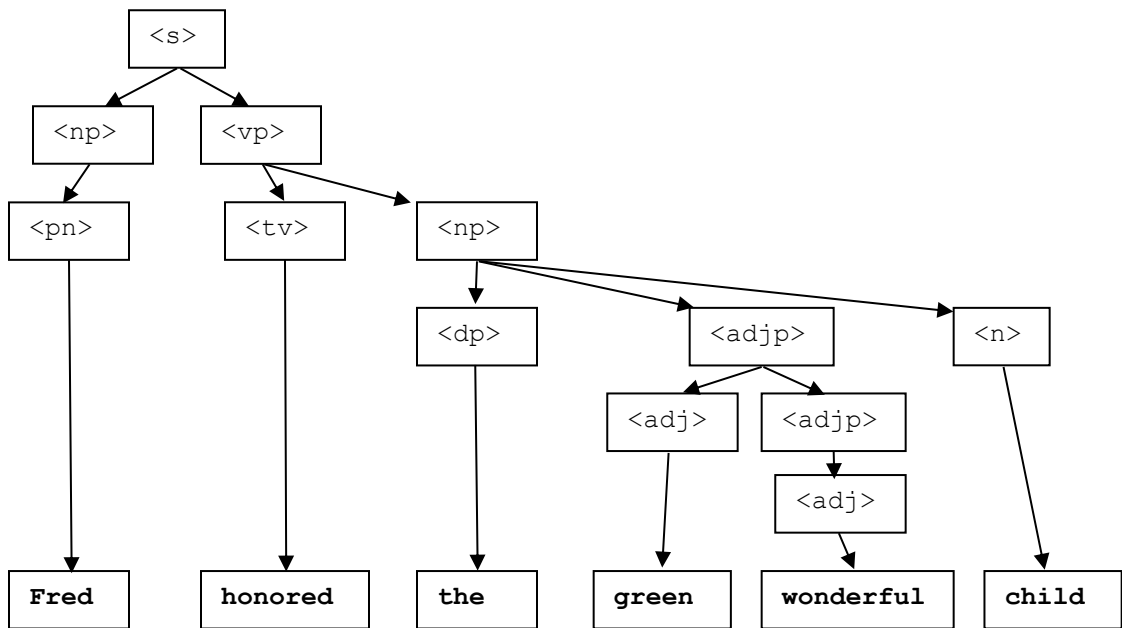
GrammarProcessor Class	
Class/Instance Data	Description
<pre>private Map<String,String[]> rules = new TreeMap<String,String[]>();</pre>	<p>The <i>nonterminals</i> of the grammar will be the key fields for the map. The rules for each nonterminal will be stored in the <code>String[]</code> array associated with the nonterminal key value. Note that since we are using a <code>TreeMap</code> implementation of the <code>Map</code> interface, the keys will be stored in sorted order.</p>
Public methods	Description
<pre>GrammarProcessor(List<String> grammar)</pre>	<p>Each entry in the <code>grammar</code> parameter holds a line from a BNF grammar file. This constructor will store the <code>grammar</code> data in the <code>rules</code> map appropriately. It should throw an <code>IllegalArgumentException</code> if the <code>grammar</code> is <code>null</code>, empty or contains more than one entry for the same nonterminal. Note that case matters, so <code><S></code> and <code><s></code> are not the same. The constructor should not change the contents of <code>grammar</code> in any way. You are strongly encouraged to display the contents of the <code>rules</code> map during development so that you are absolutely certain you have the grammar correctly stored in the map before you code the <code>generate()</code> methods described below.</p>
<pre>boolean grammarContains(String symbol)</pre>	<p>Returns true if the symbol is a nonterminal of the grammar stored in <code>rules</code>, or false otherwise.</p>
<pre>String getSymbols()</pre>	<p>Returns a <code>String</code> representation of all the nonterminal symbols of the grammar. This should be a sorted, comma-separated listing enclosed within square brackets, as in <code>"[<np>, <s>, <vp>]"</code></p> <p>Hint: Apply the <code>toString()</code> method to <code>keyset()</code> for our map.</p>
<pre>String generate(String symbol)</pre>	<p>Use the grammar in the <code>rules</code> map to randomly generate an occurrence of the given <code>symbol</code>. Throw an <code>IllegalArgumentException</code> if the grammar does not contain the given nonterminal <code>symbol</code>. <i>Call a private recursive helper method appropriately.</i></p>

GrammarProcessor Specs continued below.

GrammarProcessor Class	
Public methods	Description
String[] generate(String symbol, int times)	This method does the same thing as the preceding method, except that it returns multiple (times) randomly generated occurrences of the given symbol in a String array. It should throw an <code>IllegalArgumentException</code> under the same circumstances as above, and if times is less than 0. <i>Call a private recursive helper method appropriately.</i>

In addition to the `public generate()` methods described above, you will write `private` helper methods that perform the actual recursion. Typically, the private method has the same name as the public, but a different parameter list. To generate a random instantiation of a nonterminal, you will pick at random one of its rules and generate whatever that rule tells you to generate. This is a recursive process because a rule might lead to another nonterminal, which will lead to another nonterminal, and so on. When you finally encounter a terminal, you simply include it in the String you are generating. This becomes the base case of the recursive process. **See the `FileCrawler.java` program presented in Chapter 12 and posted in the `Lab5Files` folder. The folder also contains a dummy `crawl_dir` folder that you can ask the `FileCrawler` to display for you. Your public/private `generate()` method pairs will be very similar to the public/private `crawl()` method pairs.** Note that the recursive `crawl()` method contains a loop. There is no rule against having a loop in a recursive method.

The grammar in `sentence1.txt` can produce the sentence (`<s>`), “Fred honored the green wonderful child”, as indicated in the following diagram:



Be sure to follow the good programming style specifications posted in the Quick Links module on the Canvas website.